

Introduction(week 1, part 1)

2 Sep 2023

Basics

Logistics

- (almost) everything at the [course web page](#)
- communication/forums ([Piazza](#)); e-mail if necessary
- assignment marks (Avenue)
- Zoom/recordings (by request)

Integrity

- [notes on honesty](#)
- why copying code is good
- Stack Overflow, ChatGPT, and all that
- group work

Prerequisites

From the course outline:

- basics of linear models (as in [STATS 3A03](#)), with associated linear algebra
- basics of generalized linear models (as in [STATS 4C03/6C03](#)), including knowledge of exponential family distributions
- inferential statistics: sampling distributions, Central Limit theorem, hypothesis testing, Wald tests, maximum likelihood estimation
- ideally, *basic* knowledge of Bayesian statistics and Markov chain Monte Carlo estimation
- intermediate knowledge of R

Goals

- principles/practices of statistical modeling
 - choosing a model
 - diagnostics and troubleshooting
- good intermediate understanding of the tools (ridge/lasso, (G)(LA)MMs); unifying principles of regression modeling, shrinkage/penalized estimators
 - similarity of Bayesian and frequentist approaches
- awareness of computational foundations/scaling

Technical skills & tools

Not focal, but unavoidable and useful

- R (base + some [tidyverse](#))
- reproducibility
 - version control (Git/GitHub)
 - documents: Quarto/Sweave/Jupyter notebooks

about me

- weird background (physics/math u/g, Zoology PhD, epidemiological modeling)
- math biology (ecology/evolution/epidemiology)
- computational statistics (mixed models, Bayesian stats)

things I like/obsess about

- scientific inference \gg pure prediction (but see Navarro (2019))
- generative models
- data visualization
- solving problems in context, practical issues
- bad statistical practice (p-value abuse, snooping, dichotomania, imbalance handling, ...)

The modeling cycle

Before you start

- you need to know what the question is!
- this is hard for statisticians
- what is a large effect? what is an interesting effect?
- a low p -value is not inherently interesting!

Effect sizes

- standardized measures like Cohen's d ($(\bar{x}_1 - \bar{x}_2)/s$, where s is some measure of pooled standard deviation: [Wikipedia](#)) are common ...
- but shouldn't be used mindlessly. Real-world, unstandardized effects are usually more meaningful
- effects estimated on the log or logit scale are unitless and hence *may* be easier to generalize
- scaling predictors and responses may help (Schielzeth 2010)

An iterative process ...

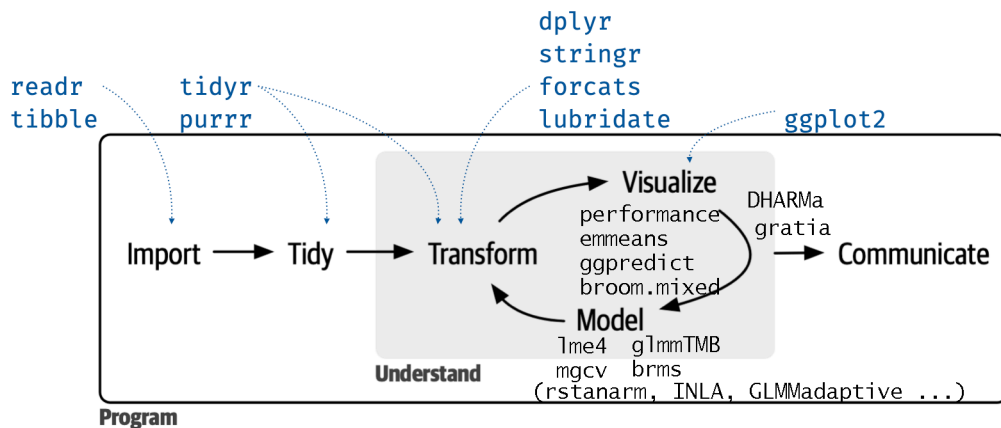


Figure 1: original from [Mine Çetinkaya-Rundel](#)



Figure 2: From Box (1976)

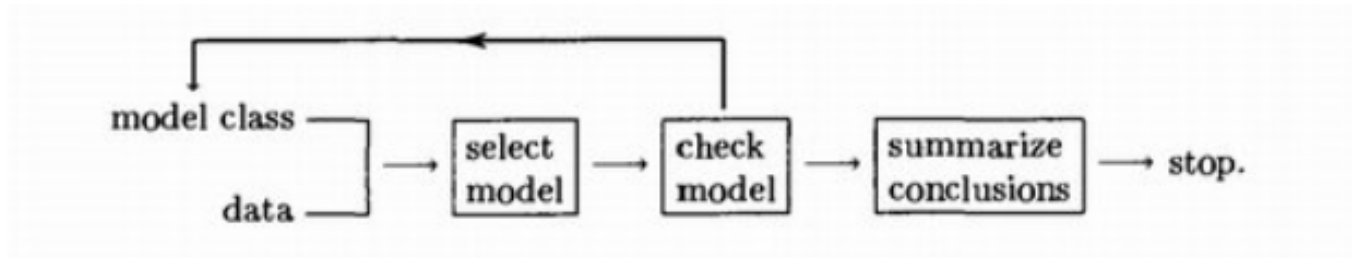


Figure 3: From McCullagh and Nelder (1989) p. 392: 'The introduction of this loop changes profoundly the process of analysis and the reliability of the final models found.'



Figure 4: from [Art Share LA](#)

Beware the garden of forking paths!

- “researcher degrees of freedom”, “HARKing”, etc.
- Simmons, Nelson, and Simonsohn (2011); Gelman and Loken (2014)

Solutions?

- pre-registration (formal or informal); report deviations from planned analysis
- choose model complexity (see Harrell RMS ch. 3), do diagnostics etc., **without reference to response variable** or metrics of significance

Choosing model complexity

- see Harrell ch. 3
- rules of thumb for inferential models with adequate power
 - e.g. $p < n/10$ or $n/20$
 - effective n depends on data type (binary < small counts < continuous)
 - how does clustering/correlation in data change effective n ?
- simplest if done **a priori**
 - data-driven choice of model complexity (e.g. by cross-validation), *while maintaining valid inference*, is delicate

Model diagnostics

- all models make assumptions
- results *may* be sensitive to **misspecification**: bias, inefficiency, inflated/deflated type I error, poor coverage ...
- hypothesis tests (e.g. Shapiro-Wilk) are deprecated

Harvey Motulsky on [CrossValidated](#):

The question normality tests answer: Is there convincing evidence of any deviation from the Gaussian ideal? With moderately large real data sets, the answer is almost always yes.

The question scientists often expect the normality test to answer: Do the data deviate enough from the Gaussian ideal to “forbid” use of a test that assumes a Gaussian distribution? Scientists often want the normality test to be the referee that decides when to abandon conventional (ANOVA, etc.) tests and instead analyze transformed data or use a rank-based nonparametric test or a resampling or bootstrap approach. For this purpose, normality tests are not very useful.

- “is there a statistically significant deviation from the model assumptions?” vs. “are the violations of the assumptions large enough to mess up my conclusions?” (**never** “are the data normally distributed?”)
- **two-stage testing** often has bad properties (H. Campbell and Dean 2014; Harlan Campbell 2021; Rochon, Gondan, and Kieser 2012; Zimmerman 2004)
- graphical diagnostics are often recommended
- but how do we judge whether deviations are too large? ([Q-Q plot survey](#))

References

- Box, George E. P. 1976. “Science and Statistics.” *Journal of the American Statistical Association* 71 (356): 791–99. <https://doi.org/10.1080/01621459.1976.10480949>.
- Campbell, Harlan. 2021. “The Consequences of Checking for Zero-Inflation and Overdispersion in the Analysis of Count Data.” *Methods in Ecology and Evolution* 12 (4): 665–80. <https://doi.org/10.1111/2041-210X.13559>.
- Campbell, H., and C. B. Dean. 2014. “The Consequences of Proportional Hazards Based Model Selection.” *Statistics in Medicine* 33 (6): 1042–56. <https://doi.org/10.1002/sim.6021>.
- Gelman, Andrew, and Eric Loken. 2014. “The Statistical Crisis in Science: Data-Dependent Analysis—a “Garden of Forking Paths”—Explains Why Many Statistically Significant Comparisons Don’t Hold Up.” *American Scientist* 102 (6): 460–60. http://link.galegroup.com/apps/doc/A389260653/AONE?u=ocul_mcmaster&sid=AONE&xid=4f4562c0.
- McCullagh, P., and J. A. Nelder. 1989. *Generalized Linear Models*. London: Chapman; Hall.
- Navarro, Danielle. 2019. “Science and Statistics.” Aarhus University. <https://slides.com/djnavarro/scienceandstatistics>.
- Rochon, Justine, Matthias Gondan, and Meinhard Kieser. 2012. “To Test or Not to Test: Preliminary Assessment of Normality When Comparing Two Independent Samples.” *BMC Medical Research Methodology* 12 (1): 81. <https://doi.org/10.1186/1471-2288-12-81>.
- Schielzeth, Holger. 2010. “Simple Means to Improve the Interpretability of Regression Coefficients: Interpretation of Regression Coefficients.” *Methods in Ecology and Evolution* 1 (2): 103–13. <https://doi.org/10.1111/j.2041-210X.2010.00012.x>.
- Simmons, Joseph P., Leif D. Nelson, and Uri Simonsohn. 2011. “False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as

Significant." *Psychological Science* 22 (11): 1359–66. <https://doi.org/10.1177/0956797611417632>.

Zimmerman, Donald W. 2004. "A Note on Preliminary Tests of Equality of Variances." *British Journal of Mathematical and Statistical Psychology* 57 (1): 173–81. <https://doi.org/10.1348/000711004849222>.