

Review of linear models

11 Sep 2023

```
## it's nice to include packages at the top
## (and NOT automatically install them)
## try not to carry over packages you don't use
library(faux)
library(brglm2) ## for lizards data
library(ggplot2); theme_set(theme_bw())
## diagnostics
library(performance)
library(DHARMA)
## downstream model evaluation
library(broom)
library(dotwhisker)
library(emmeans)
library(effects)
library(marginaleffects)
library(parameters)
## library(ggeffects)
```

Basics

- assume $\mathbf{y} \sim \text{Normal}(\mathbf{X}\beta, \sigma)$ ¹
- \mathbf{X} is the *model matrix*, can be anything we want it to be
- the *Gauss-Markov theorem* ([Wikipedia](#)) makes weaker assumptions: $\mathbf{y} = \mathbf{X}\beta + \epsilon$; as long as ϵ is mean-zero, homoscedastic with finite variance, and uncorrelated ... then the OLS solution

$$\hat{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

is the BLUE (or MVUE).

¹Notation-abuse warning ...

- we'll embrace the assumptions (which are needed for inference!)

Computation

- matrix decompositions (QR with pivoting)
- big problems: `biglm`, `speedglm`, `RcppEigen::fastLm`
 - optimized BLAS, kernel trick, etc.
 - memory vs speed vs robustness ...
 - p vs. n vs. many-small-regressions vs. ...

Inference

- σ^2 (residual variance) is $\text{RSS}/(n - p)$
- The covariance matrix is $\Sigma = \sigma^2(\mathbf{X}^\top \mathbf{X})^{-1}$.
- Individual coefficients are t -distributed
- Linear combinations of coefficients (contrasts or predictions) are t -distributed with covariance matrix $\mathbf{C}^\top \Sigma^{-1} \mathbf{C}$
- Joint hypotheses on coefficients are F -distributed
- Wald and likelihood ratio test comparisons are equivalent (but need to be careful about marginality)

Model matrices

- model definition converted to \mathbf{X} before we start
- **input variables** vs **predictor variables** (Schielzeth (2010), Gelman and Hill (2006), [CV](#))
 - transformations
 - encoding of categorical variables: **contrasts**
 - interactions
 - basis expansions (e.g. polynomials)

Wilkinson-Rogers formulas

- Wilkinson and Rogers (1973), updated by Chambers and Hastie (1991, ch. 2)
- operators: $+$, $*$, $:$, $/$, $-$, \wedge
- `I()`

Contrasts

treatment contrasts

- intercept = baseline, subsequent values are differences
- $\{\beta_0 = \mu_0, \beta_i = \mu_i - \mu_0 \text{ for } i > 0\}$
- equivalently: $\{\mu_0 = \beta_0, \mu_i = \beta_0 + \beta_i \text{ for } i > 0\}$
- **contrast matrix:**

$$\mathbf{C}\beta = \begin{pmatrix} 1 & 0 & 0 & \dots \\ 1 & 1 & 0 & \dots \\ 1 & 0 & 1 & \dots \\ \vdots & \vdots & \vdots & \ddots \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \dots \end{pmatrix} = \begin{pmatrix} \mu_1 \\ \mu_2 \\ \mu_3 \\ \dots \end{pmatrix} .$$

Maybe easier to start from the **inverse** contrast matrix: $\beta = \mathbf{C}^{-1}\mu$.

```
C <- cbind(1, contr.treatment(3)) ## R omits the intercept by default
solve(C)
```

```
      1 2 3
      1 0 0
2     -1 1 0
3     -1 0 1
```

We have to specify the baseline level (`contr.treatment` uses first level of a factor; `contr.SAS()` uses the last level).

It's nice when contrasts are *orthogonal*, i.e. all rows are independent $\rightarrow \mathbf{C}^\top \mathbf{C}$ is diagonal.

Sum-to-zero contrasts

- intercept is the (unweighted!) average rather than baseline value ($\sum \mu_i/n$)
- other parameters are differences between mean of level i and intercept ($\mu_i - \sum \mu_i/n$)
- **last** level is dropped

```
mfun <- function(C) MASS::fractions(solve(C))
(C <- cbind(1, contr.sum(3)))
```

	[,1]	[,2]	[,3]
1	1	1	0
2	1	0	1
3	1	-1	-1

```
mfun(C)
```

	1	2	3
[1,]	1/3	1/3	1/3
[2,]	2/3	-1/3	-1/3
[3,]	-1/3	2/3	-1/3

Helmert contrasts

- Weird but orthogonal
- intercept, diff of first two levels, diff of level 3 from 1 & 2, ...

```
(C <- cbind(1,contr.helmert(3)))
```

	[,1]	[,2]	[,3]
1	1	-1	-1
2	1	1	-1
3	1	0	2

```
mfun(C)
```

	1	2	3
[1,]	1/3	1/3	1/3
[2,]	-1/2	1/2	0
[3,]	-1/6	-1/6	1/3

others

- `MASS::contr.sdif()` (successive-differences)
- `contr.poly()` (orthogonal polynomial contrasts)
- custom (e.g., “none” vs “symbiont effect” vs “crabs vs shrimp” vs “two-symbiont effect”) (McKeon et al. (2012); data [here](#))

```
cc_inv <- matrix(c(1/4,1/4,1/4,1/4,
                  1,-1/3,-1/3,-1/3,
                  0,1,-1,0,
                  0,1/2,1/2,-1),
                byrow=TRUE,
                nrow=4,
                dimnames=list(c("intercept","avg_symb","C.vs.S","twosymb"),
                              c("none","C","S","CS"))
## inverse contrast matrix
MASS::fractions(cc_inv)
```

	none	C	S	CS
intercept	1/4	1/4	1/4	1/4
avg_symb	1	-1/3	-1/3	-1/3
C.vs.S	0	1	-1	0
twosymb	0	1/2	1/2	-1

```
## contrast matrix
mfun(cc_inv)
```

	intercept	avg_symb	C.vs.S	twosymb
none	1	3/4	0	0
C	1	-1/4	1/2	1/3
S	1	-1/4	-1/2	1/3
CS	1	-1/4	0	-2/3

practical issues

- too many ways to set contrasts (`options()`, `contrasts(f) <- lm(..., contrasts = list(...))`)
- terrible naming conventions: you can get used to it or use the `faux` package
- OK to fit models and later use `emmeans` to recover desired contrasts (switching linear bases)

```
mtcars$fcyl <- factor(mtcars$cyl)
lm0 <- lm(mpg ~ fcyl, mtcars)
cn <- function(x) names(coef(x))
cn(lm0)
```

```
[1] "(Intercept)" "fcyl6"          "fcyl8"
```

```
update(lm0, contrasts = list(fcyl = contr.sum(3))) |> cn()
```

```
[1] "(Intercept)" "fcyl1"          "fcyl2"
```

```
update(lm0, contrasts = list(fcyl = contr.helmert(3))) |> cn()
```

```
[1] "(Intercept)" "fcyl1"          "fcyl2"
```

using faux

```
update(lm0, data = transform(mtcars, fcyl = contr_code_sum(fcyl))) |> cn()
```

```
[1] "(Intercept)"      "fcyl.4-intercept" "fcyl.6-intercept"
```

```
update(lm0, data = transform(mtcars, fcyl = contr_code_helmert(fcyl))) |> cn()
```


```
[1] "(Intercept)" "fcyl.6-4"      "fcyl.8-4.6"
```

Interactions

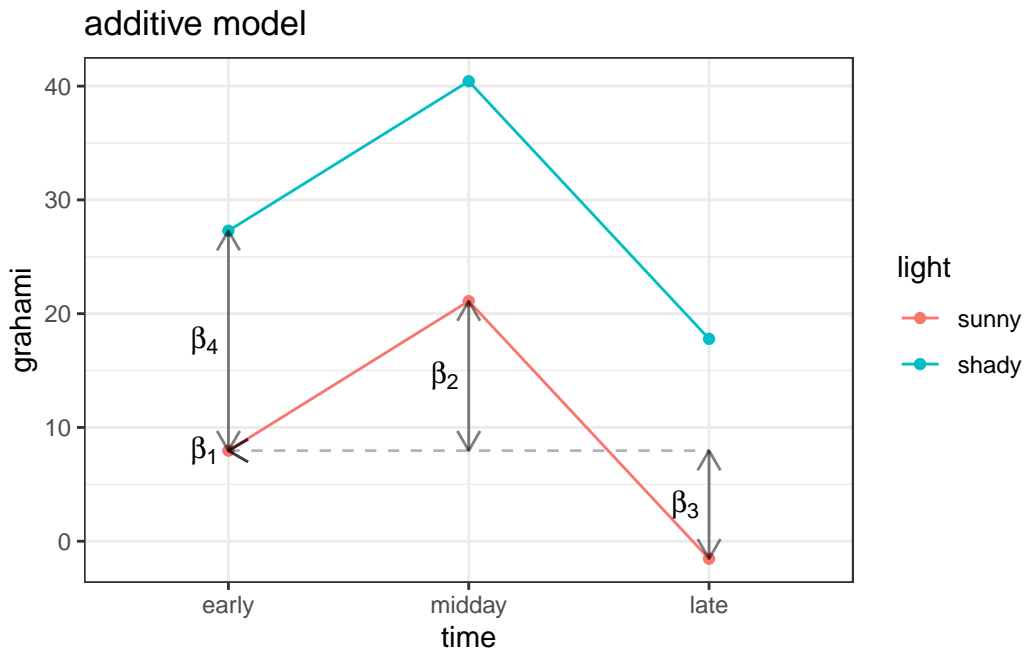
- differences in differences
- parameter values of main effects (and p values etc.) depend on contrasts/centering!
- overall model fit (R^2 , predictions, etc.) is invariant

Lizard data (Schoener (1970), from the `brglm2` package):

```
data("lizards", package = "brglm2")
ggplot(lizards, aes(time, grahami, colour = light)) +
  stat_summary(fun.data = mean_cl_boot,
               position = position_dodge(width = 0.25))
```



```
data("lizards", package = "brglm2")
lmTL1 <- lm(grahami~time+light,data=lizards)
pp <- with(lizards,expand.grid(time=levels(time),light=levels(light)))
pp$grahami <- predict(lmTL1,newdata=pp)
cc <- as.list(plyr::rename(coef(lmTL1),c(`(Intercept)`="int")))
labelpos <- with(cc,
  list(x=c(1,2,3,1),xend=c(1,2,3,1),
    y=c(int,int,int,int),
    yend=c(int,int+timemidday,int+timelate,int+lightshady)))
xpos <- -0.1
ggplot(pp,aes(x=time,y=grahami,colour=light))+geom_point()+
  geom_line(aes(group=light))+
  annotate("segment",x=labelpos$x,xend=labelpos$xend,y=labelpos$y,
    yend=labelpos$yend,alpha=0.5,
    arrow=arrow(length = unit(0.3,"cm"),ends="both"))+
  annotate("text",x=with(labelpos,(x+xend)/2)+xpos,y=with(labelpos,(y+yend)/2),
    label=paste0("beta[",1:4,"]"),parse=TRUE)+
  annotate("segment",x=labelpos$x[1],xend=labelpos$x[3],y=labelpos$y[1],
    yend=labelpos$y[1],alpha=0.3,lty=2) +
  labs(title = "additive model")
```



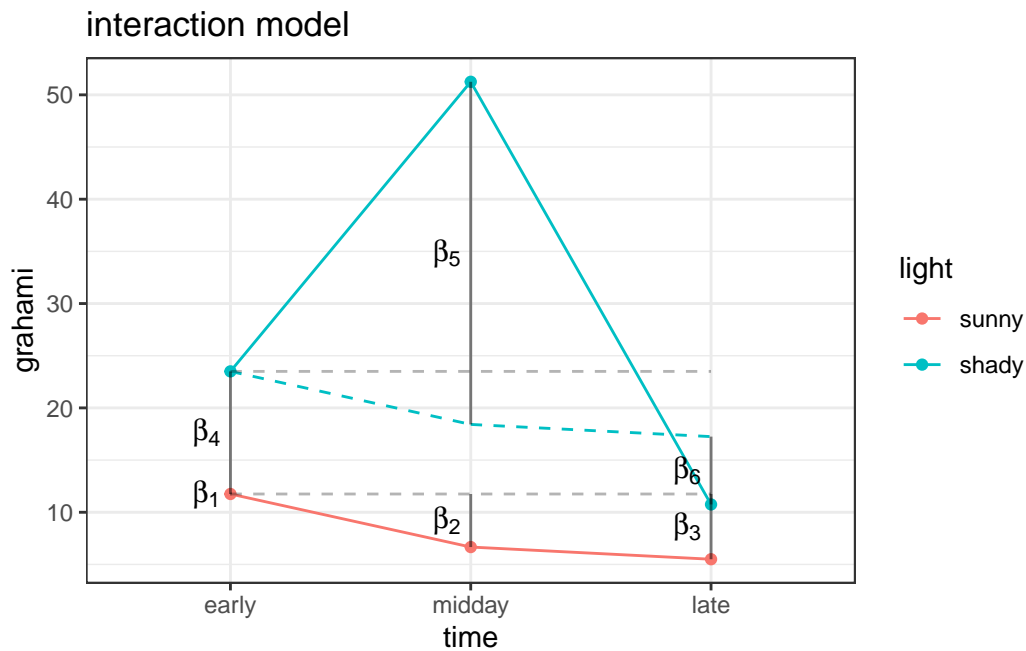
```
lmTL2 <- lm(grahami~time*light,data=lizards)
gg_color_hue <- function(n) {
  hues = seq(15, 375, length=n+1)
  hcl(h=hues, l=65, c=100)[1:n]
}
pp2 <- pp
pp2$grahami <- predict(lmTL2,newdata=pp)
cc <- as.list(plyr::rename(coef(lmTL2),c('(Intercept)\'="int",
  `timemidday:lightshady`="midshady",`timelate:lightshady`="lateshady"))))
labelpos <- with(cc,
  list(x=c(1,2,3,1,2,3),xend=c(1,2,3,1,2,3),
    y=c(int,int,int,int,int+lightshady+timemidday,int+lightshady+timelate),
    yend=c(int,int+timemidday,int+timelate,int+lightshady,
      int+timemidday+lightshady+midshady,int+timelate+lightshady+lateshady)))
xpos <- -0.1
ggplot(pp2,aes(x=time,y=grahami,colour=light))+geom_point()+
  geom_line(aes(group=light))+
  annotate("segment",x=1:2,xend=2:3,
    y=with(cc,c(int+lightshady,int+timemidday+lightshady)),
    yend=with(cc,c(int+timemidday+lightshady,int+timelate+lightshady)),
    colour=gg_color_hue(2)[2],lty=2)+
  annotate("segment",x=labelpos$x,xend=labelpos$xend,y=labelpos$y,
```



```

yend=labelpos$yend,alpha=0.5) +
  ## arrow=arrow(length = unit(0.3,"cm"),ends="both"))+
  annotate("text",x=with(labelpos,(x+xend)/2)+xpos,y=with(labelpos,(y+yend)/2),
label=paste0("beta[",1:6,"]"),parse=TRUE)+
  annotate("segment",x=rep(labelpos$x[1],2),
              xend=rep(labelpos$x[3],2),
              y=labelpos$yend[c(1,4)],
              yend=labelpos$yend[c(1,4)],alpha=0.3,lty=2) +
  labs(title = "interaction model")

```



Marginality


- Venables (1998)
- 'type (X) sums of squares'
- scaling and centering (Schielzeth 2010) alleviates many problems; sum-to-zero contrasts (weighted or unweighted?)

Model interpretation, visualization, testing

Diagnostics

- linearity,
- base R: `stats::plot.lm()`
- `performance::check_model()`
- DHARMA (`simulateResiduals(., plot = TRUE)`)

```
performance::check_model(lmTL2)
```



```
summary(lmTL2)
```

Call:

```
lm(formula = grahami ~ time * light, data = lizards)
```

Residuals:

Min	1Q	Median	3Q	Max
-30.250	-4.125	1.250	6.875	17.750

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	11.750	5.375	2.186	0.04311 *
timemidday	-5.083	8.211	-0.619	0.54408
timelate	-6.250	7.602	-0.822	0.42238
lightshady	11.750	7.602	1.546	0.14061
timemidday:lightshady	32.833	11.190	2.934	0.00927 **
timelate:lightshady	-6.500	10.751	-0.605	0.55343

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.75 on 17 degrees of freedom

Multiple R-squared: 0.7504, Adjusted R-squared: 0.677

F-statistic: 10.22 on 5 and 17 DF, p-value: 0.0001179


```
broom::tidy(lmTL2)
```

A tibble: 6 x 5

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	11.8	5.38	2.19	0.0431
2 timemidday	-5.08	8.21	-0.619	0.544
3 timelate	-6.25	7.60	-0.822	0.422
4 lightshady	11.7	7.60	1.55	0.141
5 timemidday:lightshady	32.8	11.2	2.93	0.00927
6 timelate:lightshady	-6.5	10.8	-0.605	0.553

```
## automatically drop intercept; optional by_2sd argument
```

```
dotwhisker::dwplot(lmTL2) + geom_vline(xintercept = 0, lty = 2) +  
  labs(x="difference in grahami count")
```



Interpretation and testing

- Look at coefficient tables: `summary()` or `coef(summary())`
- model comparison: `drop1()`, `anova()`, `car::Anova()`
- coefficient *plots*: `broom` + `ggplot2`, `dotwhisker`

Downstream methods

- plot predictions **with data**
- partial residuals plots (e.g. [remef package](#))
- prediction, effects plots
- uncertainty of predictions
- `emmeans`, `marginaleffects`, `effects`, `sjPlot` ...

References

- Chambers, J. M., and T. J. Hastie, eds. 1991. *Statistical Models in S*. 1st ed. Chapman & Hall/CRC.
- Gelman, Andrew, and Jennifer Hill. 2006. *Data Analysis Using Regression and Multi-level/Hierarchical Models*. Cambridge, England: Cambridge University Press.

- McKeon, C., Adrian Stier, Shelby McIlroy, and Benjamin Bolker. 2012. "Multiple Defender Effects: Synergistic Coral Defense by Mutualist Crustaceans." *Oecologia* 169 (4): 1095–1103. <https://doi.org/10.1007/s00442-012-2275-2>.
- Schielzeth, Holger. 2010. "Simple Means to Improve the Interpretability of Regression Coefficients: Interpretation of Regression Coefficients." *Methods in Ecology and Evolution* 1 (2): 103–13. <https://doi.org/10.1111/j.2041-210X.2010.00012.x>.
- Schoener, Thomas W. 1970. "Nonsynchronous Spatial Overlap of Lizards in Patchy Habitats." *Ecology* 51 (3): 408–18. <https://doi.org/10.2307/1935376>.
- Venables, W. N. 1998. "Exegeses on Linear Models." In. 1998 International S-PLUS User Conference. Washington, DC. <http://www.stats.ox.ac.uk/pub/MASS3/Exegeses.pdf>.
- Wilkinson, G. N., and C. E. Rogers. 1973. "Symbolic Description of Factorial Models for Analysis of Variance." *Journal of the Royal Statistical Society. Series C (Applied Statistics)* 22 (3): 392–99. <https://doi.org/10.2307/2346786>.