Review of linear models

25 Sep 2023

```r
## it's nice to include packages at the top
## (and NOT automatically install them)
## try not to carry over packages you don't use
library(faux)
library(brglm2) ## for lizards data
library(ggplot2); theme_set(theme_bw())
## diagnostics
library(performance)
library(DHARMa)
## downstream model evaluation
library(broom)
library(dotwhisker)
library(emmeans)
library(effects)
library(marginaleffects)
library(parameters)
## library(ggeffects)
```

Basics

- assume $\mathbf{y} \sim \text{Normal}(\mathbf{X}\beta, \sigma)$[1]
- $\mathbf{X}$ is the *model matrix*, can be anything we want it to be
- the *Gauss-Markov theorem* (Wikipedia) makes weaker assumptions: $\mathbf{y} = \mathbf{X}\beta + \epsilon$; as long as ϵ is mean-zero, homoscedastic with finite variance, and uncorrelated ... then the OLS solution
$$\hat{\beta} = (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{y}$$
  is the BLUE (or MVUE).

---

[1]Notation-abuse warning ...

- we'll embrace the assumptions (which are needed for inference!)

Computation

- matrix decompositions (QR with pivoting; see here)
- big problems: `biglm`, `speedglm`, `RcppEigen::fastLm`

  - optimized BLAS, kernel trick, etc.
  - memory vs speed vs robustness ...
  - p vs. n vs. many-small-regressions vs. ...

Inference

- σ^2 (residual variance) is $\text{RSS}/(n-p)$
- The covariance matrix is $\Sigma = \sigma^2(\mathbf{X}^\top \mathbf{X})^{-1}$.
- Individual coefficients are t-distributed
- Linear combinations of coefficients (contrasts or predictions) are t-distributed with co-variance matrix $\mathbf{C}^\top \Sigma^{-1} \mathbf{C}$
- Joint hypotheses on coefficients are F-distributed
- Wald and likelihood ratio test comparisons are equivalent
  (but need to be careful about marginality)

Model matrices

- model definition converted to **X** before we start
- **input variables** vs **predictor variables** (Schielzeth (2010), Gelman and Hill (2006), CV)

  - transformations
  - encoding of categorical variables: **contrasts**
  - interactions
  - basis expansions (e.g. polynomials)

Wilkinson-Rogers formulas

- Wilkinson and Rogers (1973), updated by Chambers and Hastie (1991, ch. 2)
- operators: `+`, `*`, `:`, `/`, `-`, `^`
- `I()`

Contrasts

treatment contrasts

- intercept = baseline, subsequent values are differences
- $\{\beta_0 = \mu_0, \beta_i = \mu_i - \mu_0 \text{ for } i > 0\}$
- equivalently: $\{\mu_0 = \beta_0, \mu_i = \beta_0 + \beta_i \text{ for } i > 0\}$
- **contrast matrix**:

$$
\mathbf{C}\beta = \begin{pmatrix} 1 & 0 & 0 & \dots \\ 1 & 1 & 0 & \dots \\ 1 & 0 & 1 & \dots \\ \vdots & \vdots & \vdots & \ddots \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \dots \end{pmatrix} = \begin{pmatrix} \mu_1 \\ \mu_2 \\ \mu_3 \\ \dots \end{pmatrix} .
$$

Maybe easier to start from the **inverse** contrast matrix: $\beta = \mathbf{C}^{-1}\mu$.

```
C <- cbind(1, contr.treatment(3))  ## R omits the intercept by default
solve(C)
```

```
  1 2 3
  1 0 0
2 -1 1 0
3 -1 0 1
```

We have to specify the baseline level (`contr.treatment` uses first level of a factor; `contr.SAS()` uses the last level).

It's nice when contrasts are *orthogonal*, i.e. all rows are independent $\rightarrow \mathbf{C}^{\top}\mathbf{C}$ is diagonal.

Sum-to-zero contrasts

- intercept is the (unweighted!) average rather than baseline value $(\sum \mu_i/n)$
- other parameters are differences between mean of level i and intercept $(\mu_i - \sum_j \mu_j/n)$
- **last** level is dropped

```
mfun <- function(C) MASS::fractions(solve(C))
(C <- cbind(1,contr.sum(3)))
```

3

```
   [,1] [,2] [,3]
1    1    1    0
2    1    0    1
3    1   -1   -1
```

```
  mfun(C)
```

```
      1    2    3
[1,]  1/3  1/3  1/3
[2,]  2/3 -1/3 -1/3
[3,] -1/3  2/3 -1/3
```

Helmert contrasts

- Weird but orthogonal
- intercept, diff of first two levels, diff of level 3 from 1 & 2, ...

```
  (C <- cbind(1,contr.helmert(3)))
```

```
   [,1] [,2] [,3]
1    1   -1   -1
2    1    1   -1
3    1    0    2
```

```
  mfun(C)
```

```
      1    2    3
[1,]  1/3  1/3  1/3
[2,] -1/2  1/2    0
[3,] -1/6 -1/6  1/3
```

others

- `MASS::contr.sdif()` (successive-differences)
- `contr.poly()` (orthogonal polynomial contrasts)
- custom (e.g., "none" vs "symbiont effect" vs "crabs vs shrimp" vs "two-symbiont effect") (McKeon et al. (2012); data here)

4

```r
cc_inv <- matrix(c(1/4,1/4,1/4,1/4,
                   1,-1/3,-1/3,-1/3,
                   0,1,-1,0,
                   0,1/2,1/2,-1),
                 byrow=TRUE,
                 nrow=4,
                 dimnames=list(c("intercept","avg_symb","C.vs.S","twosymb"),
                               c("none","C","S","CS")))
## inverse contrast matrix
MASS::fractions(cc_inv)
```

```
          none    C     S     CS
intercept  1/4  1/4   1/4   1/4
avg_symb     1 -1/3  -1/3  -1/3
C.vs.S       0    1    -1     0
twosymb      0  1/2   1/2    -1
```

```r
## contrast matrix
mfun(cc_inv)
```

```
     intercept avg_symb C.vs.S twosymb
none         1      3/4      0       0
C            1     -1/4    1/2     1/3
S            1     -1/4   -1/2     1/3
CS           1     -1/4      0    -2/3
```

practical issues

- too many ways to set contrasts (`options()`, `contrasts(f) <-`, `lm(..., contrasts = list(...))`)
- terrible naming conventions: you can get used to it or use the `faux` package
- OK to fit models and later use `emmeans` to recover desired contrasts (switching linear bases)

```r
mtcars$fcyl <- factor(mtcars$cyl)
lm0 <- lm(mpg ~ fcyl, mtcars)
cn <- function(x) names(coef(x))
cn(lm0)
```

```
[1] "(Intercept)" "fcyl6"        "fcyl8"
```

```r
  update(lm0, contrasts = list(fcyl = contr.sum(3))) |> cn()
```

```
[1] "(Intercept)" "fcyl1"        "fcyl2"
```

```r
  update(lm0, contrasts = list(fcyl = contr.helmert(3))) |> cn()
```

```
[1] "(Intercept)" "fcyl1"        "fcyl2"
```

using `faux`

```r
  update(lm0, data = transform(mtcars, fcyl = contr_code_sum(fcyl))) |> cn()
```

```
[1] "(Intercept)"      "fcyl.4-intercept" "fcyl.6-intercept"
```

```r
  update(lm0, data = transform(mtcars, fcyl = contr_code_helmert(fcyl))) |> cn()
```

```
[1] "(Intercept)" "fcyl.6-4"     "fcyl.8-4.6"
```

Interactions

- differences in differences
- parameter values of main effects (and p values etc.) depend on contrasts/centering!
- overall model fit (R^2, predictions, etc.) is invariant

Lizard data (Schoener (1970), from the `brglm2` package):

```r
data("lizards", package = "brglm2")
ggplot(lizards, aes(time, grahami, colour = light)) +
    stat_summary(fun.data = mean_cl_boot,
                 position = position_dodge(width = 0.25))
```

6

```
data("lizards", package = "brglm2")
lmTL1 <- lm(grahami~time+light,data=lizards)
pp <- with(lizards,expand.grid(time=levels(time),light=levels(light)))
pp$grahami <- predict(lmTL1,newdata=pp)
cc <- as.list(plyr::rename(coef(lmTL1),c(`(Intercept)`="int")))
labelpos <- with(cc,
  list(x=c(1,2,3,1),xend=c(1,2,3,1),
      y=c(int,int,int,int),
      yend=c(int,int+timemidday,int+timelate,int+lightshady)))
xpos <- -0.1
ggplot(pp,aes(x=time,y=grahami,colour=light))+geom_point()+
  geom_line(aes(group=light))+
  annotate("segment",x=labelpos$x,xend=labelpos$xend,y=labelpos$y,
           yend=labelpos$yend,alpha=0.5,
           arrow=arrow(length = unit(0.3,"cm"),ends="both"))+
  annotate("text",x=with(labelpos,(x+xend)/2)+xpos,y=with(labelpos,(y+yend)/2),
label=paste0("beta[",1:4,"]"),parse=TRUE)+
  annotate("segment",x=labelpos$x[1],xend=labelpos$x[3],y=labelpos$y[1],
           yend=labelpos$y[1],alpha=0.3,lty=2) +
    labs(title = "additive model")
```

additive model
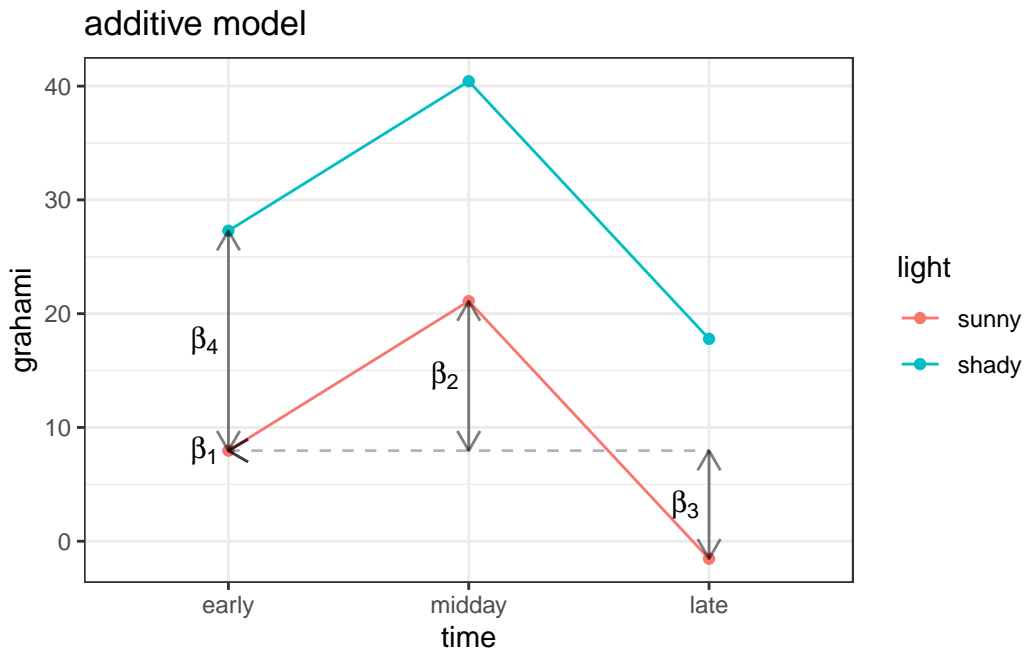
```
lmTL2 <- lm(grahami~time*light,data=lizards)
gg_color_hue <- function(n) {
  hues = seq(15, 375, length=n+1)
  hcl(h=hues, l=65, c=100)[1:n]
}
pp2 <- pp
pp2$grahami <- predict(lmTL2,newdata=pp)
cc <- as.list(plyr::rename(coef(lmTL2),c(`(Intercept)`="int",
       `timemidday:lightshady`="midshady",`timelate:lightshady`="lateshady")))
labelpos <- with(cc,
  list(x=c(1,2,3,1,2,3),xend=c(1,2,3,1,2,3),
       y=c(int,int,int,int,int+lightshady+timemidday,int+lightshady+timelate),
       yend=c(int,int+timemidday,int+timelate,int+lightshady,
              int+timemidday+lightshady+midshady,int+timelate+lightshady+lateshady)))
xpos <- -0.1
ggplot(pp2,aes(x=time,y=grahami,colour=light))+geom_point()+
  geom_line(aes(group=light))+
  annotate("segment",x=1:2,xend=2:3,
           y=with(cc,c(int+lightshady,int+timemidday+lightshady)),
           yend=with(cc,c(int+timemidday+lightshady,int+timelate+lightshady)),
           colour=gg_color_hue(2)[2],lty=2)+
  annotate("segment",x=labelpos$x,xend=labelpos$xend,y=labelpos$y,
```

8

```
              yend=labelpos$yend,alpha=0.5) +
          ## arrow=arrow(length = unit(0.3,"cm"),ends="both"))+
    annotate("text",x=with(labelpos,(x+xend)/2)+xpos,y=with(labelpos,(y+yend)/2),
  label=paste0("beta[",1:6,"]"),parse=TRUE)+
    annotate("segment",x=rep(labelpos$x[1],2),
                    xend=rep(labelpos$x[3],2),
                    y=labelpos$yend[c(1,4)],
          yend=labelpos$yend[c(1,4)],alpha=0.3,lty=2) +
      labs(title = "interaction model")
```
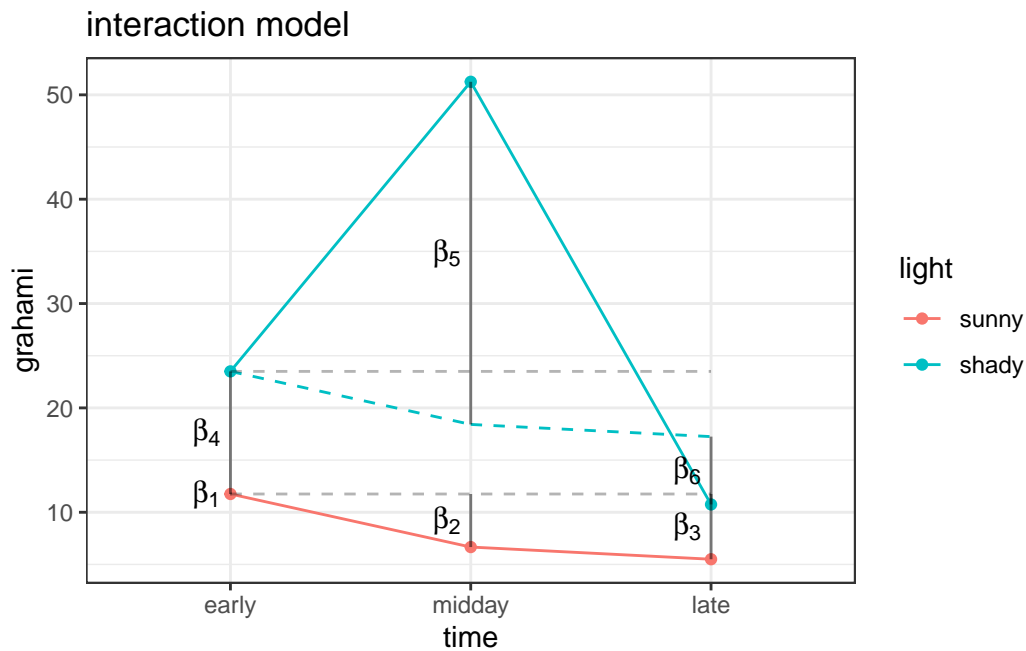

Marginality

- Venables (1998)
- 'type (X) sums of squares'
- scaling and centering (Schielzeth 2010) alleviates many problems; sum-to-zero contrasts (weighted or unweighted?)

# Model interpretation, visualization, testing

Diagnostics

- linearity > heteroscedasticity, outliers > normality
- upstream problems can induce downstream problems first
- universal plots are universal, but less interpretable than problem-specific exploration (try to identify problematic predictors/groups/etc.)
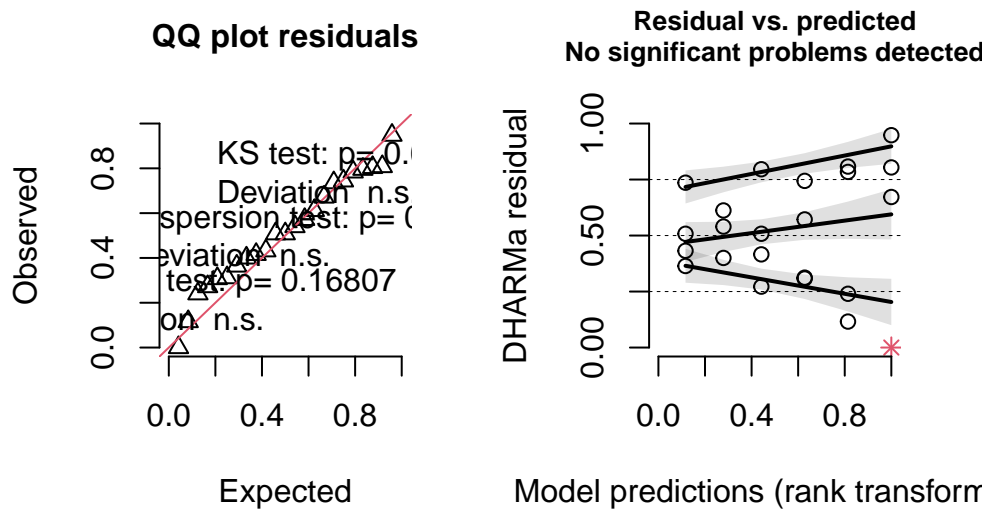
Graphical diagnostics

- base R: `stats::plot.lm()`
- `performance::check_model()`
- DHARMa (`simulateResiduals(., plot = TRUE)`)

    – (`plotResiduals(simout, form = pred_var)`)

- `broom::augment()` + plot-your-own (`ggplot2`)
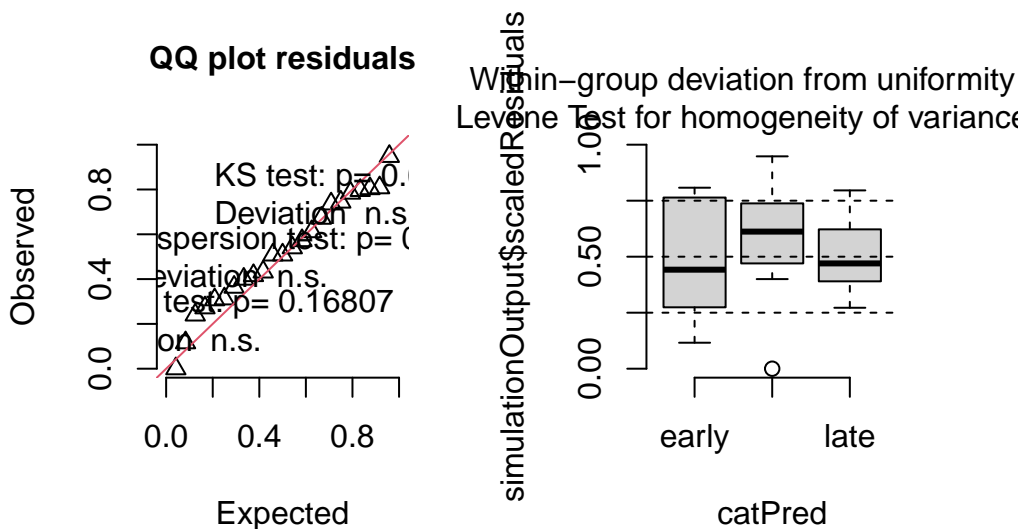
`performance::check_model(lmTL2)`

Posterior Predictive Check
Model–predicted lines should resemble obs

Linearity
Reference line should be flat and horizontal

Homogeneity of Variance
Reference line should be flat and horizonta

Influential Observations
Points should be inside the contour lines

Collinearity
High collinearity (VIF) may inflate paramete

Normality of Residuals
Dots should fall along the line

```
ss <- simulateResiduals(lmTL2)
plot(ss)
```

DHARMa residual

QQ plot residuals

Residual vs. predicted
No significant problems detected

```
plot(ss, form = lizards$time)
```

DHARMa residual

QQ plot residuals

Within−group deviation from uniformity
Levene Test for homogeneity of variance

Solutions to problems

- **nonlinearity**: transformation, add covariates (??), add interactions, add polynomial terms etc.
- **outliers**: drop values (report both!), use robust regression
- **heteroscedasticity**: transformation, model dispersion explicitly, GLMs
- **non-Normality**: transformation, GLMs

Transformation

- May do too much at once (GLMs and GAMs allow more flexibility)
- Log-transformation is often interpretable and solves problems
- Transforming boundary values (e.g. $\log(0)$) is problematic
- **Box-Cox transformations**: $y \to \frac{y^\lambda - 1}{\lambda}$ (include Jacobian term $GM^{\lambda - 1}$ in denominator to keep log-likelihood comparable)

  - flexible
  - in practice people often use 'round numbers': $\lambda = 0$ (log), $1/2$ (square root), etc.
  - `MASS::boxcox()`
  - hard to interpret!

Example

```
library(faraway)
m1 <- lm(perm ~ area, rock)
performance::check_model(m1, check = c("linearity", "homogeneity", "outliers", "qq"))
```

Linearity
Reference line should be flat and horizontal

## Homogeneity of Variance
Reference line should be flat and horizontal

## Influential Observations
Points should be inside the contour lines

Normality of Residuals
Dots should fall along the line

```r
m <- MASS::boxcox(m1, interp = FALSE)
abline(v=1/4, col = 2)
```



```r
m2 <- update(m1, perm^(1/4) ~ .)
performance::check_model(m2, check = c("linearity", "homogeneity", "outliers", "qq"))
```
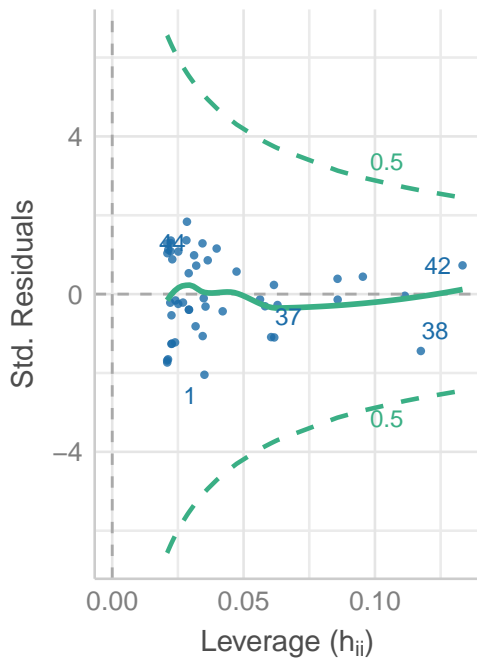
Linearity
Reference line should be flat and horizon

Homogeneity of Variance
Reference line should be flat and horizontal

Influential Observations
Points should be inside the contour lines

Normality of Residuals
Dots should fall along the line

When we transform we have to add the logs of the *Jacobian* of the transform to the log-likelihood, or subtract $\sum \log(J_i)$ from the negative log-likelihood. e.g. if we log-transform, $\sum \log(J(y_i)) = \sum \log\left(\partial(\log(y_i))/\partial y\right) = \sum \log(1/y_i) = -\sum \log y_i$. In this case $\sum \log(J(y_i)) = \sum \log((1/4)P^{(} - 3/4)) = -N\log 4 - 3/4 \sum \log(P)$:

Check:

```
logJ <- -(nrow(rock)*log(4)  + 3/4*sum(log(rock$perm)))
lm2 <- logLik(m2) + logJ
lm1 <- logLik(m1)
lm2-lm1
```

```
'log Lik.' 23.3844 (df=3)
```

```
## difference from Box-Cox sequence
max(m$y)-m$y[m$x==1]
```

```
[1] 23.52213
```

Reminder about parameter scaling

- Schielzeth (2010), Gelman (2008)
- centering ensures no problems with marginality (analogous to sum-to-zero contrasts)
- scaling predictors gives coefficients equal units (analogous to Cohen's d)
- scaling response makes coefficients *standardized partial regression coefficients* (equal to partial correlation coefficients if predictors are independent)
- scaling by 2 SD makes coefficients equivalent to those from balanced binary predictors
- scaling inputs or scaling predictors?

What about correlated predictors?

- Can compute *variance inflation factors* (VIFs)
- Dropping correlated factors is dubious: Graham (2003), Dormann et al. (2012), Morrissey and Ruxton (2018), Vanhove (2021)
- perfect collinearity gets handled automatically by R's pivoting, but may want to change contrasts/model setup

```
summary(lmTL2)
```

16

```
Call:
lm(formula = grahami ~ time * light, data = lizards)

Residuals:
    Min      1Q  Median      3Q     Max
-30.250  -4.125   1.250   6.875  17.750

Coefficients:
                        Estimate Std. Error t value Pr(>|t|)
(Intercept)               11.750      5.375   2.186  0.04311 *
timemidday                -5.083      8.211  -0.619  0.54408
timelate                  -6.250      7.602  -0.822  0.42238
lightshady                11.750      7.602   1.546  0.14061
timemidday:lightshady     32.833     11.190   2.934  0.00927 **
timelate:lightshady       -6.500     10.751  -0.605  0.55343
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.75 on 17 degrees of freedom
Multiple R-squared:  0.7504,     Adjusted R-squared:  0.677
F-statistic: 10.22 on 5 and 17 DF,  p-value: 0.0001179
```

```r
broom::tidy(lmTL2)
```

```
# A tibble: 6 x 5
  term                   estimate std.error statistic p.value
  <chr>                     <dbl>     <dbl>     <dbl>   <dbl>
1 (Intercept)                11.8      5.38      2.19  0.0431
2 timemidday                -5.08      8.21     -0.619 0.544
3 timelate                  -6.25      7.60     -0.822 0.422
4 lightshady                 11.7      7.60      1.55  0.141
5 timemidday:lightshady      32.8     11.2       2.93  0.00927
6 timelate:lightshady       -6.5      10.8      -0.605 0.553
```

```r
## automatically drop intercept; optional by_2sd argument
dotwhisker::dwplot(lmTL2) + geom_vline(xintercept = 0, lty = 2) +
    labs(x="difference in grahami count")
```

Interpretation and testing

- Look at coefficient tables: `summary()` or `coef(summary())`
- model comparison: `drop1()`, `anova()`, `car::Anova()`
- coefficient *plots*: `broom` + `ggplot2`, `dotwhisker`

Downstream methods

- plot predictions **with data**
- partial residuals plots (e.g. remef package)
- prediction, effects plots
- uncertainty of predictions
- `emmeans`, `marginaleffects`, `effects`, `sjPlot` ...

References

Chambers, J. M., and T. J. Hastie, eds. 1991. *Statistical Models in S*. 1st ed. Chapman & Hall/CRC.

Dormann, Carsten F., Jane Elith, Sven Bacher, Carsten Buchmann, Gudrun Carl, Gabriel Carré, Jaime R. García Marquéz, et al. 2012. "Collinearity: A Review of Methods to Deal with It and a Simulation Study Evaluating Their Performance." *Ecography*, no–. https://doi.org/10.1111/j.1600-0587.2012.07348.x.

Gelman, Andrew. 2008. "Scaling Regression Inputs by Dividing by Two Standard Deviations." *Statistics in Medicine* 27 (15): 2865–73. https://doi.org/10.1002/sim.3107.

Gelman, Andrew, and Jennifer Hill. 2006. *Data Analysis Using Regression and Multilevel/Hierarchical Models*. Cambridge, England: Cambridge University Press.

Graham, Michael H. 2003. "Confronting Multicollinearity in Ecological Multiple Regression." *Ecology* 84 (11): 2809–15. https://doi.org/10.1890/02-3114.

McKeon, C., Adrian Stier, Shelby McIlroy, and Benjamin Bolker. 2012. "Multiple Defender Effects: Synergistic Coral Defense by Mutualist Crustaceans." *Oecologia* 169 (4): 1095–1103. https://doi.org/10.1007/s00442-012-2275-2.

Morrissey, Michael B., and Graeme D. Ruxton. 2018. "Multiple Regression Is Not Multiple Regressions: The Meaning of Multiple Regression and the Non-Problem of Collinearity." *Philosophy, Theory, and Practice in Biology* 10 (3).

Schielzeth, Holger. 2010. "Simple Means to Improve the Interpretability of Regression Coefficients: Interpretation of Regression Coefficients." *Methods in Ecology and Evolution* 1 (2): 103–13. https://doi.org/10.1111/j.2041-210X.2010.00012.x.

Schoener, Thomas W. 1970. "Nonsynchronous Spatial Overlap of Lizards in Patchy Habitats." *Ecology* 51 (3): 408–18. https://doi.org/10.2307/1935376.

Vanhove, Jan. 2021. "Collinearity Isn't a Disease That Needs Curing." *Meta-Psychology* 5 (April). https://doi.org/10.15626/MP.2021.2548.

Venables, W. N. 1998. "Exegeses on Linear Models." In. 1998 International S-PLUS User Conference. Washington, DC. http://www.stats.ox.ac.uk/pub/MASS3/Exegeses.pdf.

Wilkinson, G. N., and C. E. Rogers. 1973. "Symbolic Description of Factorial Models for Analysis of Variance." *Journal of the Royal Statistical Society. Series C (Applied Statistics)* 22 (3): 392–99. https://doi.org/10.2307/2346786.