

MCMC examples

```
library(MCMCglmm) ## older, Gibbs-sampling
library(brms)      ## newest, lme4-like syntax, very flexible, compiled
library(rstanarm)  ## lme4-like syntax, pre-compiled
library(lme4)      ## to get data
options(brms.backend = "cmdstanr")
library(broom.mixed) ## 'tidy'
library(tidybayes)  ## convenience functions for getting MCMC output in 'tidy' format
library(bayesplot)
library(bayestestR) ## diagnostics
library(ggplot2); theme_set(theme_bw())
library(shinystan)  ## diagnostics for Stan in a Shiny window
library(tidyverse) ## general-purpose manipulations
```

- a little more on priors:

– parameter-expanded priors: $y_j | \mu, \xi_j \sim N(\mu + \alpha \sigma_j, \sigma_j^2)$, $\sigma_j \sim N(0, \sigma_\xi^2)$;
 $\alpha \sim N(\alpha_0, \sigma_\alpha)$, $\sigma_\alpha \sim \text{inverse-Gamma}(\nu)$

```
df(v/alpha.V, df1 = 1, df2 = nu, ncp = (alpha.mu^2)/alpha.V)
2 * dt(sqrt(v)/sqrt(alpha.V), df = nu, ncp = alpha.mu/sqrt(alpha.V))
```

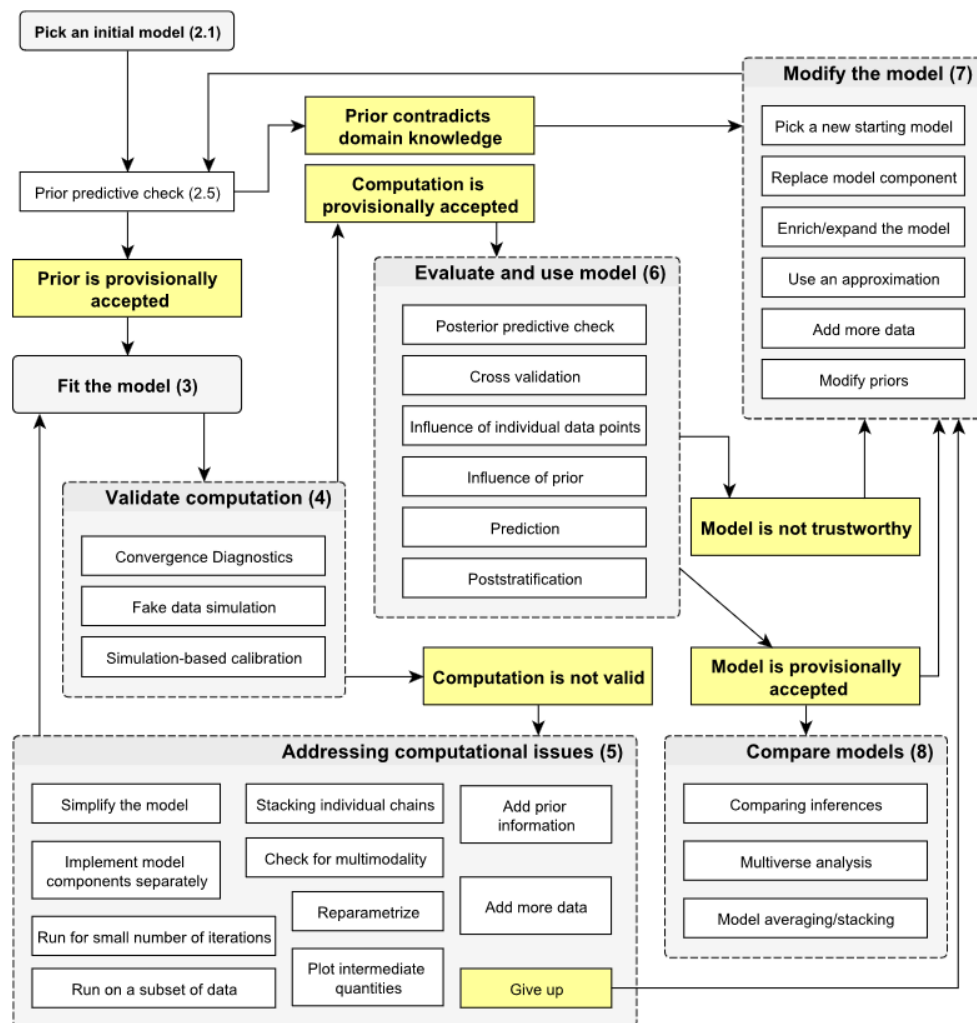
... always set `alpha.mu=0`, can set `V = 1` (or `diag()` in more complex cases) wlog;
`sqrt(alpha.V)` (scale) and `nu` are the only relevant parameters

effective sample size

- number of samples, corrected for autocorrelation
- ESS may be > sample size! (e.g. *antithetic sampling*)
- efficiency of a sampler is not (samples/time), but (effective samples/time)
- effective sample size >1000 for both tail and bulk quantities (Vehtari et al. 2021)

Bayesian workflow

Gelman et al. (2020)



simulation-based calibration

Talts et al. (2020)

default priors/prior predictive simulations:

- `rstanarm` default priors: <https://cran.r-project.org/web/packages/rstanarm/vignettes/priors.html>

Using the good old `sleepstudy` example:

```
priorpred <- stan_lmer(Reaction ~ Days + (Days|Subject),  
                      prior_PD = TRUE, data = sleepstudy, chains = 1,  
                      seed = 101,  
                      refresh = 0)
```

```
prior_summary(priorpred)
```

Priors for model 'priorpred'

Intercept (after predictors centered)

Specified prior:

~ normal(location = 299, scale = 2.5)

Adjusted prior:

~ normal(location = 299, scale = 141)

Coefficients

Specified prior:

~ normal(location = 0, scale = 2.5)

Adjusted prior:

~ normal(location = 0, scale = 49)

Auxiliary (sigma)

Specified prior:

~ exponential(rate = 1)

Adjusted prior:


~ exponential(rate = 0.018)

Covariance

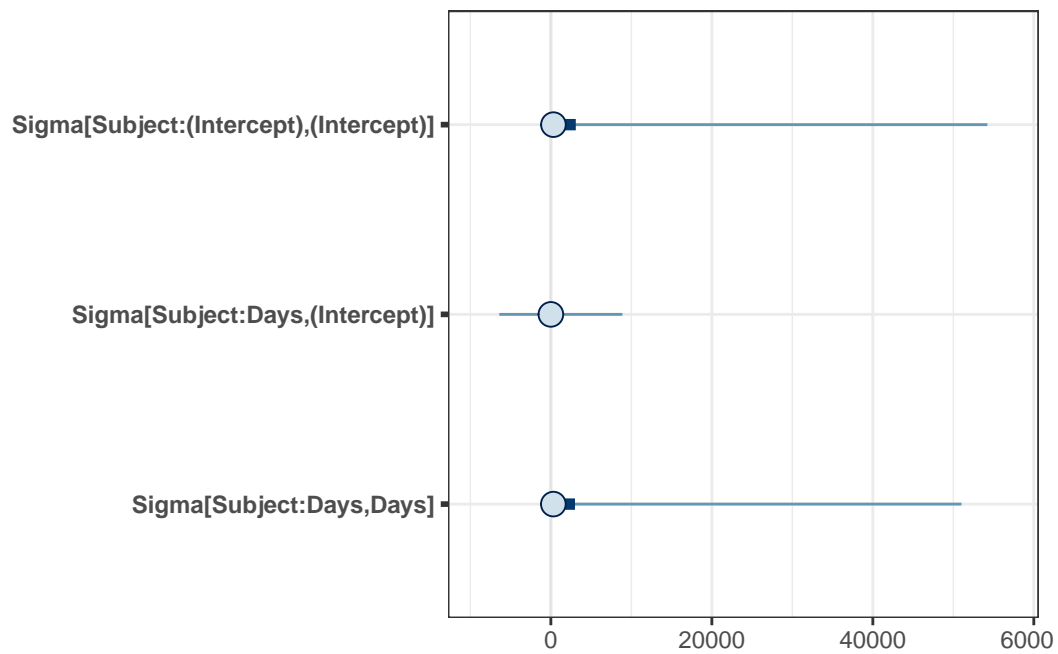
~ decov(reg. = 1, conc. = 1, shape = 1, scale = 1)

See `help('prior_summary.stanreg')` for more details

```
plot(priorpred, pars = c("(Intercept)", "Days"))
```



```
plot(priorpred, regex_pars = "Sigma")
```



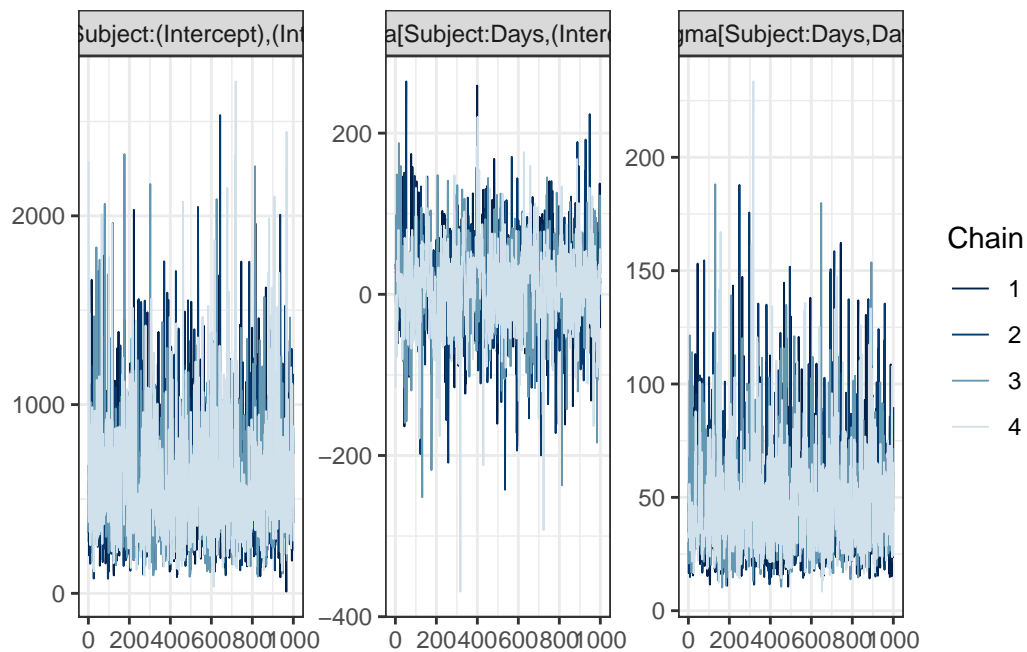
```
stanfit <- stan_lmer(Reaction ~ Days + (Days|Subject),  
  data = sleepstudy, chains = 4)
```

```
print(bayestestR::diagnostic_posterior(stanfit), digits = 4)
```


	Parameter	Rhat	ESS	MCSE
1	(Intercept)	1.001	2142	0.14337
38	Days	1.002	1353	0.04635

```
launch_shinystan(stanfit)
```

```
mcmc_trace(stanfit, regex_pars= "Sigma")
```



```
mcmc_rank_overlay(stanfit, regex_pars= "Sigma")
```



- MCMC diagnostics
 - trace plots, improved trace plots
 - R-hat Vehtari et al. (2021)
 - divergences (HMC only)

See <http://bbolker.github.io/bbmisc/bayes/examples.html>

doing stuff with the results

```
tidy(stanfit, effects=c("fixed", "ran_pars"), conf.int = TRUE)
```

```
# A tibble: 6 x 6
```

term	estimate	std.error	conf.low	conf.high	group
<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
1 (Intercept)	251.	6.35	240.	262.	<NA>
2 Days	10.4	1.63	7.48	13.1	<NA>
3 sd_(Intercept).Subject	24.2	NA	NA	NA	Subject
4 sd_Days.Subject	6.88	NA	NA	NA	Subject
5 cor_(Intercept).Days.Subject	0.0735	NA	NA	NA	Subject
6 sd_Observation.Residual	26.0	NA	NA	NA	Residual

?? why don't we get confidence intervals ?? Do it by hand ...

```
(as_draws(stanfit)
  |> tidyr::pivot_longer(everything())
  |> group_by(name)
  |> summarise(estimate = median(value),
               lwr = quantile(value, 0.025),
               upr = quantile(value, 0.975))
  |> filter(!stringr::str_detect(name, "~b\\["))
)
```

```
# A tibble: 9 x 4
  name                estimate    lwr    upr
  <chr>              <dbl>   <dbl> <dbl>
1 (Intercept)        251.    238.   264.
2 .chain              2.5      1      4
3 .draw              2000.   101.  3900.
4 .iteration           500.    26.0  975.
5 Days               10.4     6.92  13.6
6 Sigma[Subject:(Intercept),(Intercept)] 515.    161.  1376.
7 Sigma[Subject:Days,(Intercept)]        15.1 -99.4   109.
8 Sigma[Subject:Days,Days]                42.3  18.0   107.
9 sigma               25.9    23.1   29.4
```

```
form1 <- Reaction ~ Days + (Days|Subject)
get_prior(form1, sleepstudy)
```

	prior	class	coef	group	resp	dpar	nlpar	lb	ub
	(flat)	b							
	(flat)	b	Days						
	lkj(1)	cor							
	lkj(1)	cor		Subject					
student_t(3, 288.7, 59.3)		Intercept							
student_t(3, 0, 59.3)		sd							0
student_t(3, 0, 59.3)		sd		Subject					0
student_t(3, 0, 59.3)		sd	Days	Subject					0
student_t(3, 0, 59.3)		sd	Intercept	Subject					0
student_t(3, 0, 59.3)		sigma							0
source									
default									
(vectorized)									

```

      default
(vectorized)
      default
      default
(vectorized)
(vectorized)
(vectorized)
      default

```

```

b_prior <- c(set_prior("normal(200, 50)", "Intercept"),
             set_prior("normal(0, 10)", "b"),
             set_prior("normal(0, 1)", "sigma")
            )

```

```

b <- brm(form1, sleepstudy,
         prior = b_prior,
         seed = 101,           ## reproducibility
         sample_prior = 'only', ## for prior predictive sim
         chains = 1, iter = 500, ## very short sample for convenience
         silent = 2, refresh = 0 ## be vewy vewy quiet ...
        )
p_df <- sleepstudy |> tidybayes::add_predicted_draws(b)

```

‘spaghetti plot’ of prior preds

```

gg0 <- ggplot(p_df, aes(Days, .prediction, group=interaction(Subject, .draw))) +
  geom_line(alpha = 0.1)

```

```

b_prior4 <- c(set_prior("normal(200, 5)", "Intercept"),
              set_prior("normal(0, 2)", "b"),
              set_prior("normal(0, 1)", "sd"),
              set_prior("normal(0, 1)", "sigma")
             )
cc <- capture.output(
  suppressMessages(
    b_reg <- brm(form1, sleepstudy,
                prior = b_prior4,
                seed = 101,
                init = 0,
                control = list(adapt_delta = 0.95)
                ))
)

```


I've used `suppressMessages` to get rid of a lot of messages like

Chain 1 Informational Message: The current Metropolis proposal is about to be rejected because of the following issue: Exception: normal_id_glm_lpdf: Scale vector is inf, but must be positive finite! (in '/tmp/RtmpSSmixI/model-6899b70c2b466.stan', line 74, column 4 to column 55) If this warning occurs sporadically, such as for highly constrained variable types like covariance matrices, then the sampler is fine, but if this warning occurs often then your model may be either severely ill-conditioned or misspecified.

Suppressing all messages is generally a bad idea (it might suppress other messages that you do want to see), but there's no obvious way to suppress just these messages *when they occur in the warmup phase*, which seems to be a harmless case.

From the [Stan forums](#):

This is common and not a problem, the algorithm explores a large range of values in the warm-up phase and often triggers numerical problems that go away.

```
print(bayestestR::diagnostic_posterior(b_reg), digits = 4)
```

	Parameter	Rhat	ESS	MCSE
1	b_Days	0.9999	3998	0.01643
2	b_Intercept	1.0001	2952	0.05493

```
## debug(MCMCglmm::priorformat)
m <- MCMCglmm(Reaction ~ Days, random = ~us(1+Days):Subject,
             data = sleepstudy,
             verbose=FALSE,
             prior = list(G=list(G1=list(V=diag(2), nu = 0.1))))
broom.mixed::tidy(m)
```

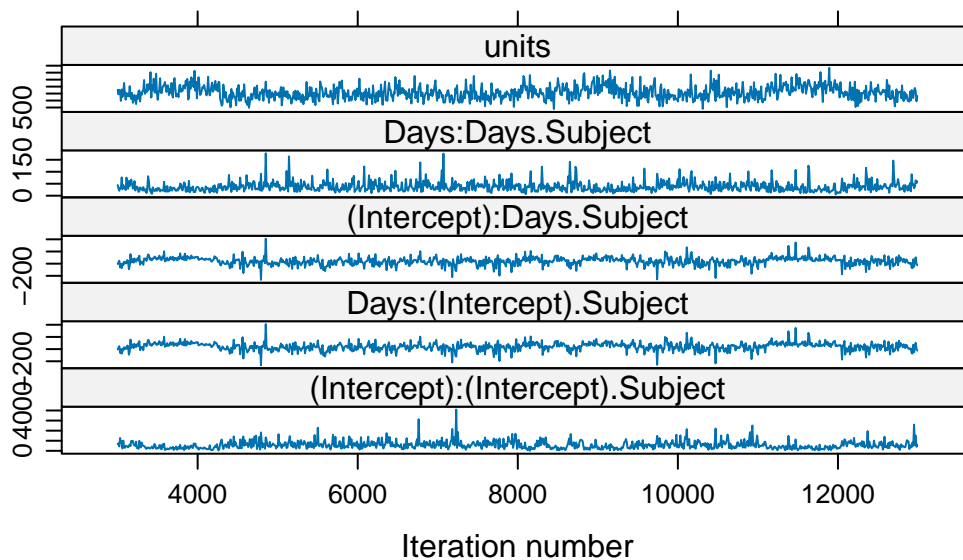
```
# A tibble: 6 x 5
  effect    group    term          estimate std.error
  <chr>    <chr>    <chr>          <dbl>    <dbl>
1 fixed   <NA>    (Intercept)    251.      6.97
2 fixed   <NA>    Days           10.5      1.63
3 ran_pars Subject var__(Intercept)  586.     384.
4 ran_pars Subject cov__(Intercept).Days  39.1     60.9
5 ran_pars Subject var__Days        37.3     17.8
6 ran_pars Residual var__Observation  705.     102.
```

```
try(MCMCglmm(Reaction ~ Days, random = ~us(1+Days):Subject,
  data = sleepstudy,
  verbose=FALSE,
  prior = list(G=list(G1=list(V=diag(2), nu = 0.1,
    alpha.mu = 0, alpha.V = diag(2))))))
```

Error in priorformat(if (NOPriorG) { :
alpha.mu is the wrong length for some prior\$G/prior\$R elements

```
m2 <- MCMCglmm(Reaction ~ Days, random = ~us(1+Days):Subject,
  data = sleepstudy,
  verbose=FALSE,
  prior = list(G=list(G1=list(V=diag(2), nu = 0.1,
    alpha.mu = rep(0,2),
    alpha.V = diag(2))))))
```

```
lattice::xyplot(m2$VCV)
```



Run longer (and thin)? Strengthen prior?

to do

- test silencing of brms messages
- improve tidy for rstanarm

- better ways to get draws
- prior pred sims for `MCMCglmm`? (examples of parameter-expansion)
- SBC examples?
- figure out compilation caching for `brms`?
- contact Hadfield about `MCMCglmm` tweaks

Gelman, Andrew, Aki Vehtari, Daniel Simpson, Charles C. Margossian, Bob Carpenter, Yuling Yao, Lauren Kennedy, Jonah Gabry, Paul-Christian Bürkner, and Martin Modrák. 2020. “Bayesian Workflow.” *arXiv:2011.01808 [Stat]*, November. <http://arxiv.org/abs/2011.01808>.

Talts, Sean, Michael Betancourt, Daniel Simpson, Aki Vehtari, and Andrew Gelman. 2020. “Validating Bayesian Inference Algorithms with Simulation-Based Calibration.” *arXiv:1804.06788 [Stat]*, October. <http://arxiv.org/abs/1804.06788>.

Vehtari, Aki, Andrew Gelman, Daniel Simpson, Bob Carpenter, and Paul-Christian Bürkner. 2021. “Rank-Normalization, Folding, and Localization: An Improved R-hat for Assessing Convergence of MCMC (with Discussion).” *Bayesian Analysis* 16 (2): 667–718. <https://doi.org/10.1214/20-BA1221>.