

# Basics of optimization/likelihood minimization

8 Oct 2024

```
library(RTMB)  
library(bbmle)
```

## References

Bolker (2008), Bolker et al. (2013)

## Minimization in statistics

- fitting models to data
- 'best fit'
- *objective function* or *loss function*; differs by application/problem type
- least squares, minimum absolute deviation, cross-entropy ...
- many of these are special cases of **negative log-likelihood**

## Maximum likelihood

- nice properties (efficient, consistent, asymptotically unbiased)
- unifying principle
- usually minimize negative log-likelihood instead
- “when it can do the job, it’s rarely the best tool for the job but it’s rarely much worse than the best” (S. Ellner)

## Minimization

- closed form solution of *score equations* (direct or via linear algebra)
- iteratively reweighted least squares
- gradient descent
- more complex iterative solutions
  - Nelder-Mead
  - quasi-Newton methods (BFGS, L-BFGS)
- automatic differentiation

## Minimization in R

- `optim()` (nlminb)
- Nelder-Mead
- various quasi-Newton methods (Press et al. 2007)

```
X <- model.matrix(~wool*tension, data=warpbreaks)
y <- warpbreaks$breaks
nll <- function(beta) {
  mu <- exp(X %*% beta)
  -sum(dpois(y, mu, log=TRUE))
}
par0 <- rep(0, ncol(X))
opt1 <- optim(par0, nll, control = list(maxit = 1000))
par2 <- coef(glm(breaks~wool*tension, data=warpbreaks, family=poisson))
all.equal(opt1$par, unname(par2), tolerance = 1e-3)
```

[1] TRUE

## with mle2

- wrapper for `optim`, variant of `stats4::mle()`
- data argument
- better accessor methods (`coef()`, `vcov()`, `broom::tidy()`, `profile()`, `confint()`, `predict()`, ...)

```
names(par0) <- parnames(nll) <- paste0("beta", 1:ncol(X))
confint(mle2(nll, par0))
```

	2.5 %	97.5 %
beta1	3.69723021	3.8930351
beta2	-0.61480841	-0.3002960
beta3	-0.78557346	-0.4545355
beta4	-0.76139971	-0.4328077
beta5	0.39936387	0.8783493
beta6	-0.06653957	0.4428584

## formula notation

```
mle2(breaks~dpois(exp(eta)),
      parameters = list(eta ~ wool*tension),
      start = list(eta = 0),
      data = warpbreaks)
```

Call:

```
mle2(minuslogl = breaks ~ dpois(exp(eta)), start = list(eta = 0),
      data = warpbreaks, parameters = list(eta ~ wool * tension))
```

Coefficients:

eta.(Intercept)	eta.woolB	eta.tensionM	eta.tensionH
3.7966810	-0.4565402	-0.6186293	-0.5957747
eta.woolB:tensionM	eta.woolB:tensionH		
0.6381522	0.1882803		

Log-likelihood: -228.48

## mle2 notes

- uses BFGS (with finite difference gradients!) by default; may want Nelder-Mead
- probably want to provide a link function for dispersion parameters, e.g. fit negative binomial with `size = exp(logsize)`

## Template Model Builder and autodiff

- everything is better with gradients
- **finite difference** gradients (R default) are terrible (expensive and inaccurate)

- FD vs symbolic vs automatic/algorithm
- magic/chain rule; *cheap gradient* principle (Kristensen et al. 2016)

## deriv() in base R

Useful, illustrative, but limited

```
deriv(expression(cos(x*sin(x^2))), "x")
```

```
expression({
  .expr1 <- x^2
  .expr2 <- sin(.expr1)
  .expr3 <- x * .expr2
  .value <- cos(.expr3)
  .grad <- array(0, c(length(.value), 1L), list(NULL, c("x")))
  .grad[, "x"] <- -(sin(.expr3) * (.expr2 + x * (cos(.expr1) *
    (2 * x))))
  attr(.value, "gradient") <- .grad
  .value
})
```

## RTMB package

```
data("prussian", package = "pscl")
prussmin <- prussian[1:3,]
X2 <- model.matrix(~ factor(year), data = prussmin)
par2 <- list(beta = rep(0, ncol(X2)))
f2 <- function(pars) {
  getAll(pars) ## this is like with() or attach()
  mu <- exp(X2 %*% beta)
  -sum(dpois(prussmin$y, lambda = mu, log = TRUE))
}
```

